

SRE Book Builder — Operator Guide

Turns a folder of scanned pages into one PDF that can be searched, and that has a clickable list of articles down the side.

Print this and keep it by the workstation.

Before you start

You need one folder containing the scans for **one issue**, and nothing else. The files must be TIFFs, one page per file, named so they sort in reading order — ..._01.tif, ..._02.tif, and so on. The first file must be the front cover.

The program never changes your scans. It only reads them.

The window

The screenshot shows the SRE Book Builder application window. The 'Issue' tab is selected, displaying the following information:

- Issue folder: C:\Users\milesm\Documents\repos\Snake River Echoes\Image Files\SRE Vol 1 Number 3
- 24 sheets found.
- Title: Snake River Echoes
- Volume: 1, Issue: 3, Year: 1971
- Publisher: Upper Snake River Valley Historical Society
- Sheet: 1 is printed page 49
- Image quality: Balanced (200 DPI)

Below the issue information is a 'Bookmarks' section with a table listing the contents of the issue:

Title	Printed page	Sheet
Front Cover	49	1
Table of Contents	50	2
A Critical Look at Historical Editing	51	3
The Massacre on Birch Creek	52	4
Idaho Merchants Tokens	56	8
Idaho Tokens	58	10
Annual Fall Public Meeting	60	12
Dedication of Marker #378	64	16
Daughters of Utah Pioneers	64	16
I'm Going Back Home	65	17
Long Distance Rides and Raids	66	18
Library History Collections	70	22
Western History Books	71	23
Upper Snake River Valley Historical Society	72	24
Board of Directors	72	24

At the bottom of the window, there is a 'Confirm' button, an 'Add' button, a 'Remove' button, and a 'Build PDF' button. A status bar at the bottom indicates '15 bookmarks proposed. 7 to keep or remove.'

The whole window, part way through reviewing an issue.

Everything happens here, top to bottom: choose the folder, check the issue details, review the bookmarks, build.


Doing an issue

1. Open the program — *SRE Book Builder* in the Start menu.

2. Click Browse and pick the issue's folder.

Reading the text starts straight away and takes about a minute for a 20-page issue. You can fill in the issue details while it works.

3. Check the issue details.



Most of this is filled in from the file names.

This picture predates the "Pages missing from this scan" box, which now sits below the page-number row.

Correct anything wrong. Two settings are worth a second look:

Sheet [1] is printed page [49]

This tells the finished PDF what page numbers to show. The program works it out from the numbers printed in the corners of the pages. **Check it against the issue in your hand.** See *Sheets and printed pages* below.

Pages missing from this scan

Usually empty. If it lists page numbers, the folder is short of those pages — see *If pages are missing from the scan* below. You can correct the list, and everything after it rennumbers as you type.

4. Review the bookmarks. This is the part only a person can do.

Click any bookmark to see that page on the right. Drag the divider between the list and the page to make the page bigger, or maximise the window.

5. Click Build PDF.

It takes a second or two, because the text was already read in step 2. The PDF appears in an output folder inside the issue's folder.

If you change a bookmark afterwards, just press Build again — it is quick.

Bookmarks in red

Bookmarks		
Title	Printed page	Sheet
Front Cover	49	1
Table of Contents	50	2
A Critical Look at Historical Editing	51	3
The Massacre on Birch Creek	52	4
Idaho Merchants Tokens	56	8
Idaho Tokens	58	10
Annual Fall Public Meeting	60	12
Dedication of Marker #378	64	16
Daughters of Utah Pioneers	64	16
I'm Going Back Home	65	17
Long Distance Rides and Raids	66	18
Library History Collections	70	22
Western History Books	71	23
Upper Snake River Valley Historical Society	72	24
Board of Directors	72	24

Black bookmarks are settled. Red ones want a decision from you.

Black entries are settled: the program found the article's heading on that page and is confident. **Red entries want a decision.** There are three kinds, and they need different things from you.

The status line at the bottom of the window counts them separately — for example *"2 to keep or remove, 1 needs a page number"* — so you can see at a glance which kind of work is left.

"This is printed here, but the contents page never listed it"

The program found a heading printed on a page — a poem, a book list, the board of directors — that the issue's own contents page does not mention. The page number is right. The only question is whether it belongs in the outline.

Keep it with **Confirm**, or drop it with **Remove**.

"The contents page named this, but I could not find where it starts"

The contents page lists an article, but the program could not find its heading in the body — usually because the heading is worded differently, or the page is a photograph it cannot read. These show a — instead of a page number, because the number it has is only a placeholder.

Give it a page number, then confirm it.

These sit **at the bottom of the list**, below everything with a real page, because the program has no idea where they go. Do not go looking for them in page order.

"This one is waiting on a page that was never scanned"

The contents page lists an article starting on a page the folder does not contain. This is not yours to fix, and the program does not ask you to place it: there is no page to point it at until someone rescans.

It shows the printed page it is waiting for, and stays in your saved review. It is left out of the PDF. Scan the page, build again, and it takes its place. See *If pages are missing from the scan* below.

The buttons



The buttons under the bookmark list.

Button	What it does
Confirm	Accepts the selected bookmark as it stands
Add	A new bookmark, for something the program missed entirely
Remove	Deletes the selected bookmark
↑ ↓	Moves a bookmark up or down the list
-> Indent	Tucks a bookmark under the one above, for poems inside a poetry section
<- Outdent	Brings it back out to the top level

The program will not build until every red bookmark is dealt with. That is deliberate: a bookmark that jumps to the wrong page is worse than no bookmark, because a reader has no way to tell.

Finding the right page

Use the issue's own contents page. It gives a **printed page number** — the number printed on the paper. Type that straight into the **Printed page** column and the program works out which scan it is.

Contents page says the article is on page **56**. Double-click the **Printed page** cell and type **56**. The Sheet column fills itself in.

You never need to do that arithmetic yourself. If you type a page the issue does not have, it will say so rather than accept it.

Then click the bookmark and check the picture on the right really does show that article starting. That check takes two seconds and is the whole point of the preview.

Sheets and printed pages

These are different, and the difference matters.

- **Sheet** — which scan it is. Sheet 1 is the first file, always.

- **Printed page** — the number printed on the paper in 1971.

They are often not the same. *Snake River Echoes* was paginated continuously through each volume, so an issue can start at printed page 27, or 49, or 73. The cover and contents pages usually carry no printed number, but they are still counted.

If sheet 1 is printed page 49, the finished PDF shows 49 (1 of 24) — so a citation to page 56 takes a reader to page 56.

When the numbering starts after the cover

In many issues the cover and the contents page carry no printed number at all, and printed page 1 is the third sheet. Say so with the same setting:

Sheet [3] is printed page [1]

The two sheets before it are then numbered **i** and **ii**, the way a book numbers its front matter, and the Printed page column shows those roman numerals so you can see it worked. You can type **ii** back into that column just as you would type 14.

Do not try to give the cover a printed page of **0**. There is no page 0, and the program will refuse it. Front matter is what that setting is for.

If every bookmark suddenly points past the end of the issue, you have almost certainly typed printed page numbers into the **Sheet** column. Use the **Printed page** column instead. The program will usually spot this and say so.

If pages are missing from the scan

Sometimes the program will warn you, right after it reads an issue:

The contents page refers to printed page(s) 28, 29, which no scan in this folder carries.

It works this out two ways. It reads the numbers printed on the pages themselves — if one scan prints 27 and the very next prints 30, then 28 and 29 are not in the folder. And it compares the contents page's own citations against the pages it can see, which catches pages missing off the end of an issue, where there is no later number to show the jump.

This is worth stopping for. Check the issue against the paper copy.

It is not a guess about your work. Of the first five issues run through this program, three turned out to be missing pages that nobody had noticed.

What the program does about it

The numbering steps over the gap. If sheet 3 is page 27 and sheet 4 is page 30, that is what the finished PDF says. The pages keep the numbers printed on them; nothing is renumbered to close the hole.

The article on the missing page keeps its bookmark. It sits in the list marked as waiting on a page that is not in this scan, showing the page number it belongs on. It is **not** written into the PDF — a bookmark that jumps to the wrong page is worse than none, because a reader cannot tell. It stays in your saved review, so when the pages are scanned and you build again, it comes back in the right place on its own.

You have the last word on the list. The **Pages missing from this scan** box in the issue panel is yours to edit. The program is reading OCR of a page printed in 1971 and can misread a number, so if it has this wrong, correct it — everything after renumbers as you type.

It will not build until you answer. Either correct the list, or click **I have checked these**. That is deliberate: publishing a book that is quietly short of two pages is the thing this whole check exists to prevent.

Getting the missing pages in later

Scan the pages that were missed, put them in the folder in reading order, and open the issue again. Click **Re-analyse** so the program reads the new scans, check the page numbers, and build. The bookmark that was waiting takes its place.

The other settings

Image quality — leave on *Balanced (200 DPI)* unless you have a reason. *High* makes a much larger file; *Small* is for emailing. Changing it does not mean re-reading the text, so it is cheap to try.

Re-analyse — reads the scans again and throws away the bookmarks currently listed, including any you edited. Use it if you open an issue that was done some time ago and the list looks wrong or empty. It asks before discarding anything.

What you end up with

Inside the issue folder, an output folder containing:

File	What it is
..._Vol1_No3.pdf	The finished book. This is the one to keep and share
..._Vol1_No3.srebook.json	Your review, saved. Keeps the PDF rebuildable later
.cache	Working files. Large, and safe to delete

The PDF opens showing the bookmark list. Readers can search it with Ctrl-F.

If something goes wrong

"There are no page scans (.tif) in ..." The folder has no TIFFs, or you picked the folder above the right one.

"Could not read the scan . _SOMETHING.TIF" You are on a version before 0.1.4. Those "._" files are left behind by Macs writing to a shared drive; they are not scans. Update, and they are ignored.

"... holds 24 pages in one file." The scans are one big TIFF rather than one file per page. Split it and try again — the program needs a file per page so it can bookmark and number them.

An article from the contents page cannot be found anywhere. Check whether its page is in the scan at all. A missing page is the commonest reason, and the program warns about it when it can tell.

The bookmark list is nearly empty. Some issues set their contents page in a way the program reads poorly. Add the bookmarks by hand with **Add**, using the contents page. The PDF is still fully searchable either way.

The cover and contents page are being given printed numbers, and should not be. Set **Sheet [3] is printed page [1]** — or whichever sheet is really printed page 1. See *When the numbering starts after the cover* above.

The list looks like an old, worse result. Opening an issue shows the review you saved last time. Click **Re-analyse** to read the scans afresh.

A page looks crooked or blank in the preview. Check the original scan. The program straightens pages automatically but cannot fix a bad scan. A page it cannot read is still included — it just has no searchable text on it.

It refuses to build and lists problems. Each line names one bookmark and what it needs. Deal with them and press **Build** again.

What this program will not do

It does not check whether the bookmark titles are *correct* — only that each one points at a real page. If the contents page has a typo, so will the bookmark, unless you fix it.

You are the accuracy check. The program is a head start, not an authority.